

Constraint Programming

on Integers — a practical tour with NuCS

Yan Georget · 60 minutes

What you'll get out of this hour

Part 1 History & theory

Where CP comes from, NP-completeness, what a CSP is.

Part 2 Real problems

Knapsack, TSP, job-shop — modeled in NuCS.

Part 3 Propagation

The reasoning engine, with diagrams.

Part 4 Search & heuristics

Backtracking, branching choices, branch-and-bound.

Part 5 Inside NuCS


How a fast solver is actually built.

No prior CP knowledge assumed. We stay on **integer variables** throughout, and use **NuCS** for concrete code.

Part 1

History & theory

A short history of CP



CP grew out of the idea: **state the constraints, let the machine search.** Forty years of making that search smart.

Why we need CP: NP-completeness

A problem is in **NP** if a candidate solution can be *checked* in polynomial time.

A problem is **NP-hard** if every NP problem reduces to it. **NP-complete** = $\text{NP} \cap \text{NP-hard}$.

Classic NP-complete problems you've already met:

- **SAT** — is this Boolean formula satisfiable?
- **Graph coloring** — can this graph be k -colored?
- **Hamiltonian cycle** — a tour visiting every node once?
- **Bin packing, subset sum, scheduling with precedences, ...**

No known polynomial algorithm — yet we need answers every day. CP is one principled way to attack them.

The naive approach blows up fast

8-queens — place 8 queens on a board, none attacking another.

Encoding	Search space
Naive: pick any 8 squares	$\binom{64}{8} \approx 4.4 \cdot 10^9$
One queen per column	$8^8 \approx 1.7 \cdot 10^7$
Permutations (one per column and row)	$8! = 40\,320$
A good CP model	~100 nodes explored

The art is **encoding** the problem so the solver can prune aggressively — not brute force.

A CP problem in one breath

A **Constraint Satisfaction Problem (CSP)** is a triple $\langle X, D, C \rangle$:

- $X = \{x_1, \dots, x_n\}$ – **variables**
- $D = \{D_1, \dots, D_n\}$ – a **domain** per variable (here: finite sets of integers)
- $C = \{c_1, \dots, c_m\}$ – **constraints**, each restricting a subset of variables

A **solution** assigns each x_i a value $v_i \in D_i$ so that every $c \in C$ holds.

Add an objective to **minimize/maximize** and it becomes a **COP** (Constraint Optimization Problem).

Where CP sits among solvers

Paradigm	Variables	Sweet spot
SAT	Boolean	Huge formulas, conflict learning
MIP	Continuous + integer	Linear models, strong LP relaxation
SMT	Boolean + theories	Verification, symbolic execution
CP	Integer, finite	Structured, non-linear combinatorics

CP shines with **global constraints**: `all_different`, `circuit`, `cumulative`, `element`, `table`, ...
— each one a reusable chunk of combinatorial reasoning.

Teaser: N-Queens in NuCS

```
from nucs.problems.problem import Problem
from nucs.propagators.propagators import ALG_ALLDIFFERENT
from nucs.solvers.backtrack_solver import BacktrackSolver

class QueensProblem(Problem):
    def __init__(self, n: int):
        super().__init__([(0, n - 1)] * n) # n vars, domain 0..n-1
        self.add_propagator(ALG_ALLDIFFERENT, range(n)) # columns
        self.add_propagator(ALG_ALLDIFFERENT, range(n), range(n)) # diagonal

        self.add_propagator(ALG_ALLDIFFERENT, range(n), range(0, -n, -1)) # diagonal


print(next(BacktrackSolver(QueensProblem(8)).solve()))
```

Three constraints. The diagonals reuse `all_different` with a per-variable **offset** — no extra variables.

Part 2

Real problems, modeled in NuCS

The modeling recipe



The skill is in **B** and **C**: a good choice of variables makes the constraints natural and the propagation strong.

We'll see this three times: **knapsack**, **TSP**, **job-shop**.

Problem 1 — Knapsack

Items with **weight** and **volume**; a bag of fixed **capacity**. Pick a subset that fits and **maximizes total weight**.

Model:

- one **0/1 variable** x_i per item.
- $\sum_i \text{volume}_i x_i \leq \text{capacity}$.
- maximize $\sum_i \text{weight}_i x_i$.



Knapsack in NuCS

```
class KnapsackProblem(Problem):
    def __init__(self, dataset):
        n = len(weights)
        super().__init__([(0, 1)] * n + [(0, sum(weights))]) #  $x_i \in \{0,1\}$ , + total
        #  $\sum_i \text{volume}_i \cdot x_i \leq \text{capacity}$ 
        self.add_propagator(ALG_LINEAR_LEQ_C, range(n), [*volumes, capacity])
        #  $\sum_i \text{weight}_i \cdot x_i = \text{total}$ 
        self.add_propagator(ALG_LINEAR_EQ_C, range(n + 1), [*weights, -1, 0])

solver.maximize(problem.weight)
```

0/1 variables + two linear constraints. The objective variable `total` is **part of the model** — that's what makes optimization prune.

Problem 2 — Travelling Salesman

n cities, cost c_{ij} to go from i to j . Find a tour visiting each city once, of **minimum total cost**.

Successor model — the key trick:

- succ_i = the city right after i .
- **circuit** forces *one* Hamiltonian cycle — rules out **sub-tours**.
- $\text{cost}_i = c_{i, \text{succ}_i}$ via **element** (index a row by a variable).



"No sub-tours" and "indexing by a variable" are exactly the kind of reasoning CP packages as **global constraints**.

TSP in NuCS

```
class TSPProblem(CircuitProblem):
    def __init__(self, costs):
        n = len(costs)
        super().__init__(n) # succ/pred vars + circuit + no-sub-cycle
        self.succ_costs = self.add_variables(...) # one cost var per city
        self.total_cost = self.add_variable(...)
        for i in range(n):
            # cost_i = costs[i][succ_i] ← element: index a row by a variable
            self.add_propagator(ALG_ELEMENT_EQ, [i, self.succ_costs + i], costs[i])
        # total_cost = sum of per-city costs
        self.add_propagator(ALG_SUM_EQ,
                           list(range(self.succ_costs, self.succ_costs + n)) + [self.total_cost])

solver.minimize(problem.total_cost)
```

NuCS also keeps a redundant **predecessor** model — same tour seen from both ends doubles the propagation.

Problem 3 — Job-shop scheduling


n jobs, m machines. Each job is a chain of operations; each runs on a given machine for a fixed time.

A machine

does **one operation at a time**. Minimize the **makespan** (finish time of the last operation).

Variables: one **start time** per operation.

- **Precedence:** op k before $k+1 \rightarrow$ a linear inequality.
- **Resource:** one **disjunctive** per machine — no overlap.
- **Makespan** = max of the job completions.



Job-shop in NuCS

```
class JobShopProblem(Problem):
    def __init__(self, jobs):
        ... # one start var per operation
        for j in range(self.job_nb):
            for k in range(self.machine_nb - 1): # job precedence
                self.add_propagator(ALG_LINEAR_LEQ_C,
                                    [self.start(j, k), self.start(j, k + 1)], [1, -1, -durations[j][k]])
        # makespan = max of completion times
        self.add_propagator(ALG_MAX_EQ, completions + [self.makespan])
        for machine in range(self.machine_nb): # one resource constraint per machine
            self.add_propagator(ALG_DISJUNCTIVE, machine_starts, machine_durations)
```

The disjunctive propagator does serious work: overload checking, **edge-finding**, detectable precedences. When demands aren't unit, the same idea generalizes to **cumulative** (a resource of capacity > 1).

Lessons from the three models

1. **Variable choice is the lever** — 0/1 for knapsack, successors for TSP, start-times for scheduling.
2. **Few high-level constraints** beat many primitive ones (`circuit` , `disjunctive` , linear).
3. **The objective variable lives in the model** — propagation tightens it, branch-and-bound exploits it.
4. **Redundant constraints** (TSP predecessors, scheduling cumulative) strengthen reasoning for free.

A CP model is mostly **declarative**: *what* must hold, not *how* to search.

Part 3

Propagation — the reasoning engine

Two engines, one loop

Solving a CSP interleaves:

1. **Propagation** — each constraint removes values that *cannot* be in any solution. Repeat to a fixpoint.
2. **Search** — when propagation stalls, **branch** (guess); on failure, **backtrack**.

Propagation alone

rarely finishes a problem.

Search alone

explodes exponentially.

Together they're powerful: propagation shrinks the tree that search must explore.

Domains as intervals (bound consistency)

NuCS represents each domain as an **interval** $[\text{min}, \text{max}]$ — two integers, not a full set.

```
super().__init__([(0, n - 1)] * n) # n domains, each = interval [0, n-1]
```

- **Cheap:** 2 ints per variable, mutated in place.
- **Ground** when `min = max`.
- **Trade-off:** no "holes" inside an interval (that's *arc* consistency) — but far faster per call.

Bound consistency + strong global constraints is the modern sweet spot, and what NuCS uses.

Propagation by example: $x + y = 10$


Domains: $x \in [0, 8]$, $y \in [3, 7]$.

Step	Reasoning	New domains
start	—	$x \in [0, 8]$, $y \in [3, 7]$
filter x	$x = 10 - y$, so $x \geq 10 - 7$ and $x \leq 10 - 3$	$x \in [3, 7]$
filter y	$y = 10 - x$, so $y \geq 10 - 7$ and $y \leq 10 - 3$	$y \in [3, 7]$

Each constraint pushes the **bounds** inward using the others' bounds. No guessing yet.

A propagator is just: *read min/max, tighten min/max, report what changed.*

The propagation fixpoint



Only re-wake constraints whose variables actually changed. Terminates because domains are finite and only shrink.

Global constraints: more than sugar

`all_different(x_1, ..., x_n)` is *logically* a clique of $\binom{n}{2}$ pairwise `≠`.

Clique of `≠`

propagates **weakly** — each `≠` only fires once a side is ground.

Dedicated `all_different`

reasons over the **whole set**: Hall intervals (BC), bipartite matching (Régin's AC).

Example: if $\{x_1, x_2, x_3\}$ all have domain $[1, 3]$, they *use up* $\{1, 2, 3\}$ — so any other variable can drop those three values. Pairwise `≠` can't see this.


Same logical meaning, **much more pruning per call** → dramatically smaller search tree.

Part 4

Search & heuristics

Backtracking search

When propagation stalls, pick a variable, split its domain, recurse — undo on failure.



A **trail** stores the bounds before each branch, so backtracking is just restoring two integers per touched variable.

Search in pseudocode

```
def search(state):  
    propagate(state) # fixpoint  
    if inconsistent: return FAIL  
    if all variables ground: yield SOLUTION; return  
    x = choose_variable(state) # variable heuristic  
    for v in choose_values(D_x): # value heuristic  
        push(state);  
        assign(x, v) # record on the trail  
        search(state)  
        pop(state) # backtrack: restore bounds
```

Two heuristics decide the *shape* of the tree — and the shape is everything.

Heuristics: choosing where to branch

Variable heuristics

- **first-fail**: smallest domain first → hit dead-ends early.
- most-constrained / max-degree.
- impact- / activity-based.

Value heuristics

- smallest value first.
- value that *least* constrains others.
- problem-specific (e.g. "schedule earliest").

NuCS ships both, and lets you **stage** them: e.g. job-shop branches on a *critical machine* first, scheduling each task as early as possible — orders of magnitude fewer nodes than a generic order.

Good model + good heuristic = the whole ball game.

Optimization = repeated satisfaction

To minimize an objective variable z :

```
best = +∞
while True:
    add
    temporary
    constraint
    z < best
    sol = solve() # find one solution
    if sol is None: return best # proven optimal
    best = sol[z] # tighten, keep going
```

Each solution adds a tighter bound on z , which **propagates** and prunes the rest of the tree.


This is **branch-and-bound**. In NuCS: `solver.minimize(var)` / `solver.maximize(var)`.

Because z is a normal variable, the bound $z < \text{best}$ is just another constraint the engine already knows how to exploit.

Part 5

Inside NuCS — how a fast solver is built

Anatomy of the solver



Everything is a **NumPy array** — domains, propagator tables, the queue — so the hot loop can be **JIT-compiled with Numba**.

A propagator is three functions

Registered once, identified by an integer `ALG_*`:

```
ALG_X = register_propagator(get_triggers_x, get_complexity_x, compute_domains_x)
```

- `get_triggers_x(...)` → `EVENT_MASK` — which events of *this* variable should re-wake me? (`MIN`, `MAX`, `GROUND`)
- `get_complexity_x(...)` → `int` — used to order the queue (cheap propagators first).
- `compute_domains_x(domains, parameters)` → `status` — the filtering itself; returns `INCONSISTENCY` / `CONSISTENCY` / `ENTAILMENT`.

Entailment = "this constraint can never be violated again" → skip it for the rest of the subtree.

A real propagator: $x \leq y + c$

```
@jit(cache=True, fastmath=True)
def compute_domains_leq_c(domains, parameters):
    x, y = domains[0], domains[1]
    c = parameters[0]
    if x[MAX] <= y[MIN] + c:          # already true for every value → entailed
        return PROP_ENTAILMENT
    x[MAX] = min(x[MAX], y[MAX] + c)  # x can be at most y's max (+ c)
    if x[MIN] > x[MAX]: return PROP_INCONSISTENCY
    y[MIN] = max(y[MIN], x[MIN] - c)  # y must be at least x's min (- c)
    if y[MIN] > y[MAX]: return PROP_INCONSISTENCY
    return PROP_CONSISTENCY
```

Pure mutation of `min / max`. No allocation, no Python objects, no branching on types — exactly what Numba compiles to tight machine code.

Events & triggers — re-wake only what matters

```
def get_triggers_leq_c(n, variable, parameters):  
    if variable == 0:          #  $x \leq y + c$ : only x's MIN can tighten y  
        return EVENT_MASK_MIN  
    return EVENT_MASK_MAX      # only y's MAX can tighten x
```

When `compute_domains` moves a bound, the engine looks up `triggers[variable, event]` and wakes exactly those propagators.

Wakeup discipline is half the performance story — never recompute a constraint nothing touched.

Idempotency: returning at a fixpoint

A subtle but important contract: when a propagator returns, calling it again on the same domains should change **nothing**.

- The engine treats "*queue empty*" as "*global fixpoint reached*" — valid only if each propagator is **idempotent on return**.
- Some rules aren't naturally idempotent (e.g. scheduling edge-finding): one pass leaves more pruning on the table.
- → such propagators run an **internal loop** to their own fixpoint before returning.

Cheaper than letting the global queue re-schedule the propagator over and over.

Why this is hard to make fast in Python

- Per-variable work is **tight integer math** — compares, mins, maxes.
- A solve can issue **millions** of propagator calls per second.
- Plain CPython interpreter overhead would dominate by 50–100×

NuCS strategy:

1. **Numba JIT** every hot function — `@njit(cache=True, fastmath=True)`.
2. **No Python objects** in jitted code — only typed NumPy arrays and ints.
3. **Mutate in place**, never reallocate.
4. **Function pointers** so propagators can be dispatched from inside jitted code.
5. **Multiprocessing** over `problem.split(...)` for embarrassingly parallel search.

What you give up vs. what you get



You give up

- Sparse domains with **holes** (NuCS is bound-consistent).
- Arbitrary Python in hot paths.
- Some expressiveness — model with the propagators that exist (or write one).



You get

- A **pure-Python** package competitive with C++ solvers on many benchmarks.
- A **small, readable** codebase you can extend in an afternoon.
- Warm-cache (Numba) startup + trivial multiprocessing.

When should you reach for CP?



Good fit

- **Scheduling, rostering**, configuration, routing, puzzles.
- Rich logical structure — `all_different`, precedences, packing.
- Integer objective, highly constrained model.



Less good fit

- Pure linear/continuous → **MIP**.
- Boolean-heavy, little structure → **SAT**.
- Very-large-scale industrial OR → **CP-SAT** (OR-Tools) blends CP + SAT.

What we covered

1. **History & theory** — CP states constraints and searches smartly; NP-completeness is why.
2. **Modeling** — TSP, knapsack, job-shop: the *variable choice* is the lever.
3. **Propagation** — bound consistency to a fixpoint; global constraints carry the reasoning.
4. **Search** — backtracking + heuristics + branch-and-bound.
5. **NuCS** — a propagation queue + a backtracking loop, made fast with **Numba**.

Resources



- **CSPLIB** — csplib.org — problem catalogue (Golomb #6, knapsack #133).



- **Handbook of Constraint Programming** — Rossi, van Beek, Walsh (Elsevier, 2006).



- **MiniZinc Handbook** — the best on-ramp to modeling.



- **NuCS** — github.com/yangeorget/nucs



- **OR-Tools CP-SAT** — Google's industrial-grade solver.

Questions?

Bonus demo, if there's time:

live-run Queens, knapsack, job-shop (mt06), TSP-15.